

14<sup>th</sup> Feb. 2022  
Monday.

## Unit - 2.

### Transaction management & concurrency control

#### Transaction.

- \* Defi. → A transaction is any action that reads from all writes to a database.
- A transaction can be consist of SELECT, UPDATE and INSERT statement.

#### A. Transaction property. [Acid property]

- A transaction consist of following properties:
  - Atomicity
  - Consistency
  - Isolation
  - Durability
  - Serializability.

#### 1) Atomicity.

- It requires that all operations of a transaction must be completed.
- If a transaction  $T_1$  has  $n$  SQL request, all  $n$  request must be successfully completed otherwise the entire transaction is aborted.

## 2) Consistency.

- It indicates the permanence of database consistent state.
- When a trans. is completed, the database must be in a consistent state.

## 3) Isolation.

It means that the data use during the execution of a trans. can not be used by a second trans. until the first one is completed.

## 4) Durability.

It ensures that once trans. changes are done and committed, they can not be undone or lost even in the event of a system failure.

## 5) Serializability

It ensures that the schedule for the concurrent execution of a trans. displays consistent result.

## B. Transaction log.

- DBMS uses a transaction log to keep track of all trans. that update the database.
- The DBMS uses the information stored in this log for a recovery requirement.



- The trans. log consist of following information:
- A record for beginning of the trans.
  - SQL statement
  - Name of the object affected by the trans. [tablename]
  - The before and after value
  - The pointers to the previous and next trans. log entries
  - The ending [commit] of the trans.

Trans ID	Trans no.	PREV PTR	NEXT PTR	Operation	Table	Attribute	Before value	After value
1	101	null	2	START	---			
2	101	1	3	UPDATE	PRODUCT	Qty	35	135
3	101	2	4	UPDATE	CUSTOMER	Balance	1000	90
4	101	3	null	COMMIT	---			

\* Concurrency control.

- The parallel execution of a trans. in multi-user database system is called concurrency control.
- Its objective is to ensure the serializability of trans. in a multi-user database environment.
- Concurrency control has 3 main problems:
- Lost update
  - Uncommitted data
  - Inconsistent retrieval

## A) Lost update

→ The lost update problem occurs when 2 trans.  $T_1$  and  $T_2$  are updating the same data element and one of the update is lost.

e.g.,

Two concurrent trans.  $T_1$  and  $T_2$  update QTY value for some item in the PRODUCT table.

$T_1 \rightarrow$  Purchase 100 unit

$T_2 \rightarrow$  Sell 30 unit

Following table shows serial execution of trans. and provides correct result:

Time	Trans.	Step	Stored value
1	$T_1$	Read Qty	85
2	$T_1$	Qty $\rightarrow$ Qty + 100	
3	$T_1$	Write Qty	135
4	$T_2$	Read Qty	135
5	$T_2$	Qty $\rightarrow$ Qty - 30	
6	$T_2$	Write Qty	105

Following table shows the problem of lost update:

- The first trans.  $T_1$  has not committed when trans.  $T_2$  is executed.
- The value of trans.  $T_1$  is overwritten by trans.  $T_2$ .



Time	Trans.	Step	Stored value
1	$T_1$	Read QTY	35
2	$T_2$	Read QTY	35
3	$T_1$	$QTY = 35 + 100$	
4	$T_2$	$QTY = 35 - 30$	
5	$T_1$	Write QTY	135
6	$T_2$	Write QTY	5

### B) Uncommitted data.

→ The problem of uncommitted data occurs when 2 trans.  $T_1$  and  $T_2$  are executed concurrently and  $T_1$  is rollback after  $T_2$  has already accessed the uncommitted data.

→ It violates the isolation property of trans.

e.g.,

The 2 concurrent trans.  $T_1$  and  $T_2$  updates QTY in the PRODUCT table.

$T_1 \rightarrow$  Purchase 100 unit

$T_2 \rightarrow$  Sell 30 unit

Following table shows serial execution of this trans. and displays correct answer:

Time	Trans.	Step	Stored value
1	$T_1$	Read Qty	35
2	$T_1$	$Qty \rightarrow Qty + 100$	
3	$T_1$	Write Qty	135
4	$T_1$	Roll back	35
5	$T_2$	Read Qty	35
6	$T_2$	$Qty = Qty - 30$	
7	$T_2$	Write Qty	5

Following table shows the problem of uncommitted data when rollback is completed:

Time	Trans.	Step	Stored value
1	$T_1$	Read Qty	35
2	$T_1$	$Qty = Qty + 100$	
3	$T_1$	Write Qty	135
4	$T_2$	Read Qty	135
5	$T_2$	$Qty = Qty - 30$	
6	$T_1$	Roll back	35
7	$T_2$	Write Qty	105



15<sup>th</sup> Feb. 2022  
Tuesday.

classmate  
Date \_\_\_\_\_  
Page 69

### c) Inconsistent retrievals.

- It occurs when a transaction accesses data before and after one or more other transactions finish working with such data.
- The problem is that the transaction might read some data before they are changed and other data after they are changed which results in inconsistent retrievals.

### D) The scheduler.

- It is a special DBMS process that establishes the order in which the operations are executed within a concurrent transaction.
- The scheduler provides execution of database operations to ensure isolation and serializability of transactions.
- Its main job is to create a serializable schedule of a transaction operation.
- It also makes sure that the computer's CPU and storage system are used efficiently.
- The problem with that approach is that the processing time is wasted when the CPU waits for a read or write operation to finish.
- In the following table, 4 operations are in conflict when they access the same data and at least one of them is a write operation.

Operations	Transaction		Result
	T <sub>1</sub>	T <sub>2</sub>	
Read	Read	Read	No conflict
Read	Read	Write	Conflict
Write	Read	Read	Conflict
Write	Write	Write	Conflict

→ Several methods are used like locking, time stamping, optimistic methods.

\* Concurrency control with locking methods.

- A lock guarantees exclusive use of data item to a current trans.
- It means that trans. ~~T<sub>2</sub>~~ T<sub>2</sub> does not have access to a data item that is currently used by trans. T<sub>1</sub>.
- The use of lock based on the assumption that conflict between trans. is likely which is called pessimistic methods [locking].
- All lock information is handled by the lock manager which is responsible for assigning the locks used by the trans.

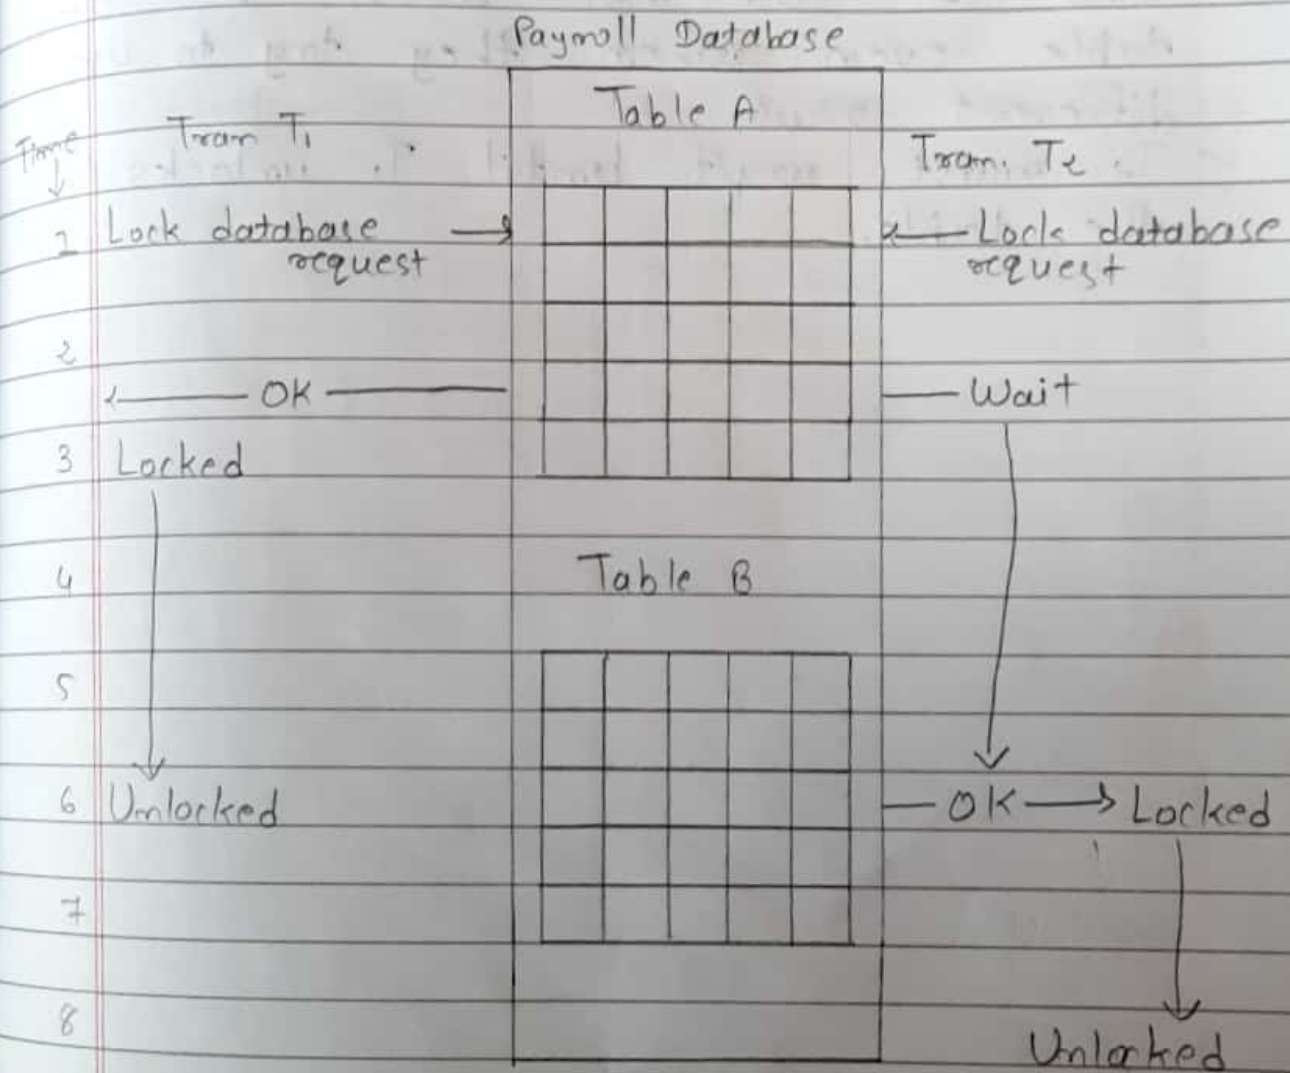
A) Lock granularity/level.

- It indicates the level of lock used.
- Locking can be take place at the database, table, page, row or even field [attribute].



Q.1.) Database level.

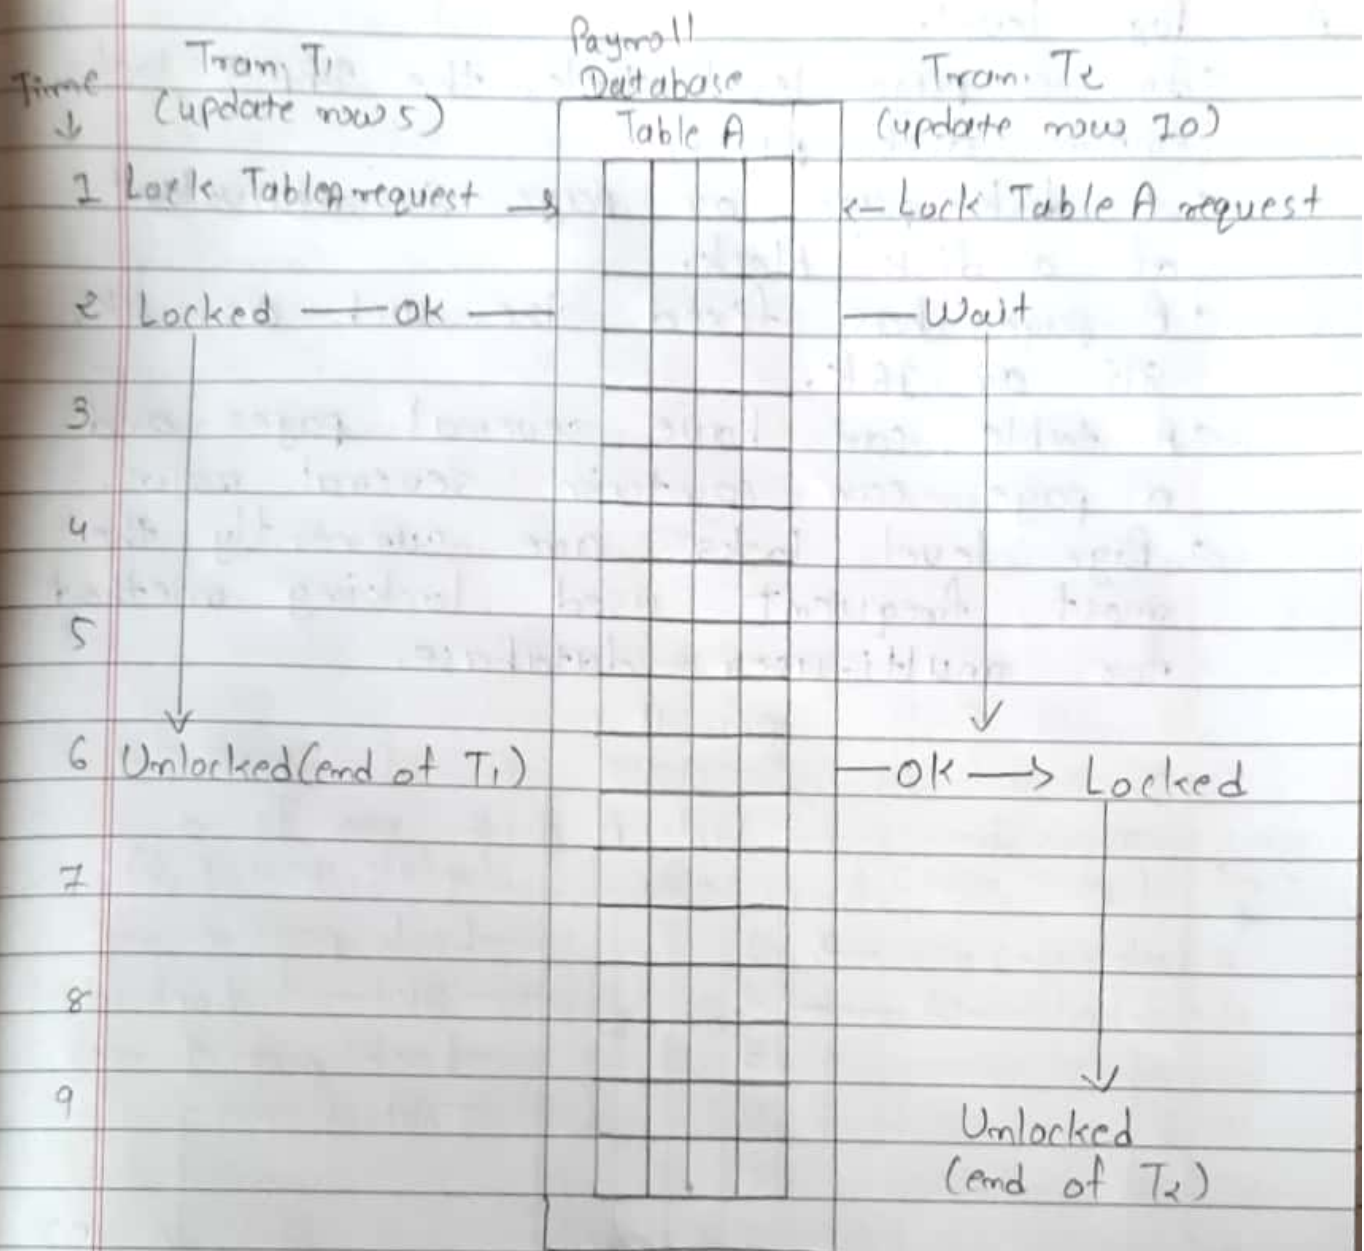
- In a database level lock, the entire database is locked.
- It prevents the use of any tables in the database by trans.  $T_2$  while trans.  $T_1$  is executed.
- The following figure shows database level lock in which trans.  $T_1$  and  $T_2$  can not access the same database concurrently even when they use different tables.



A.2.) Table level.

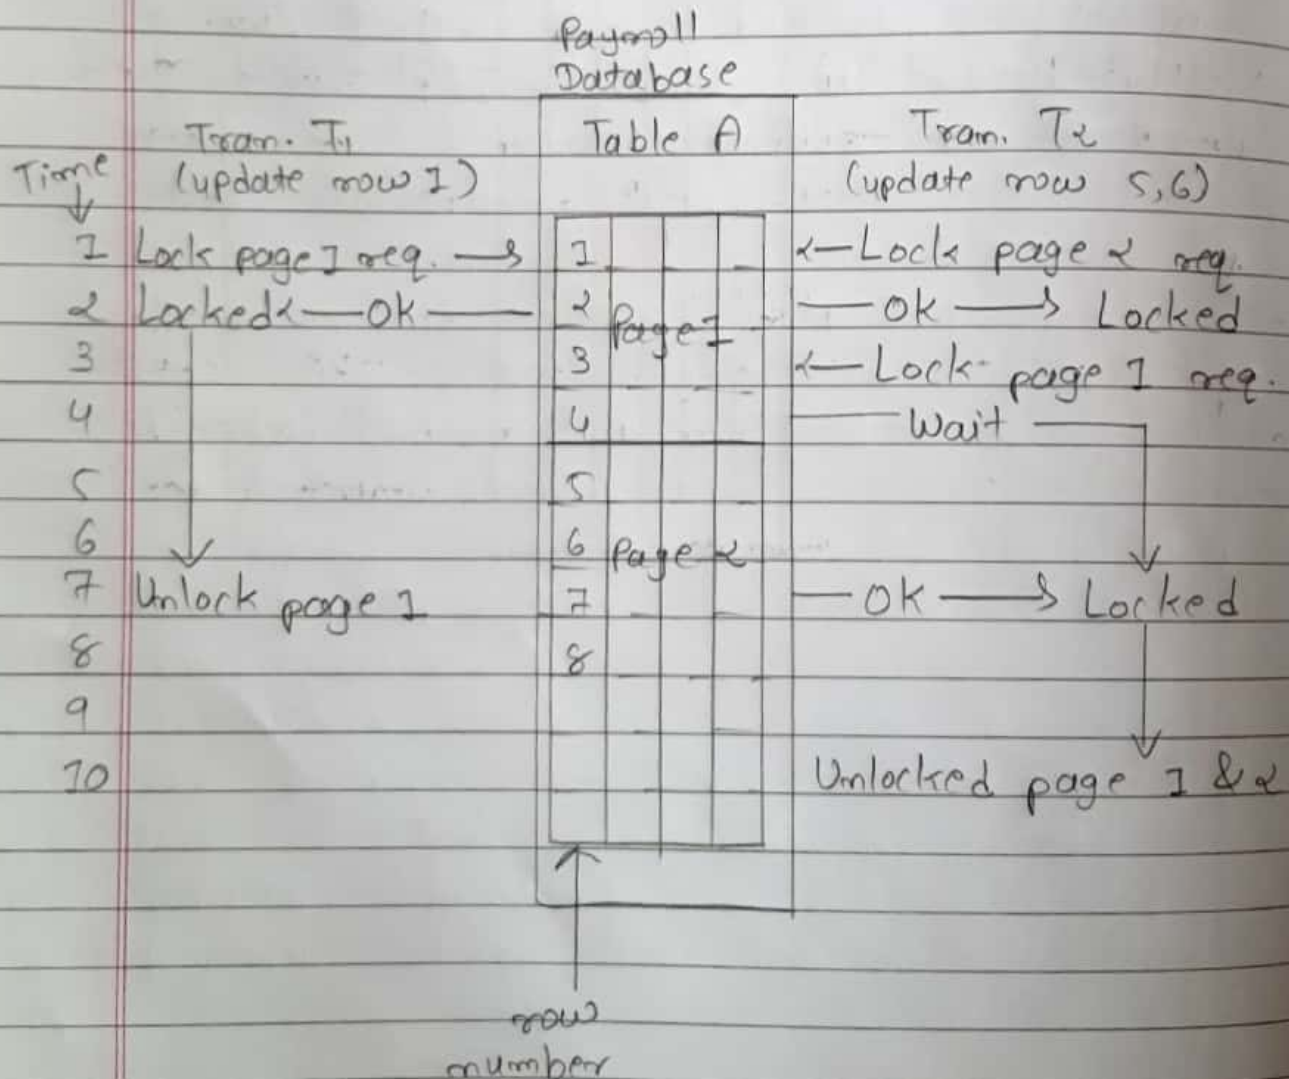
- In a table level lock, the entire table is locked.
- It prevents the access of any row by trans.  $T_2$  while trans.  $T_1$  is using the table.
- If a trans. requires access to several tables, each table maybe locked.
- Table level locks are not suitable for multi-user database.
- In the following figure trans.  $T_1$  and  $T_2$  can not access the same table even when they try to use different rows.
- $T_2$  must wait until  $T_1$  unlocks the table.





### A.3.) Page level.

- In a page level lock, the DBMS locks entire disk page.
- A disk page or page is equivalent of a disk block.
- A page has fixed size such as 4K, 8K or 16K.
- A table can have several pages and a page can contain several rows.
- Page level locks are currently the most frequent used locking method for multi-user database.





21<sup>st</sup> Feb. 2022  
Monday.

classmate

Date  
Page

75

Rolling level

It is much less restrictive than the above locks.

The dbms allows trans. to access different rows of the same table even when the rows are locked on the same page.

The following example shows both trans. can execute certainly even when the requested rows are on the same page.

Payroll database

Trans. T <sub>1</sub> (update row 1)	Table A	Trans. T <sub>2</sub> (update row 2)
1 Lock row 1 req. →	1	← Lock row 2 req.
2	2 Page 1	
3	3	
4	4	
5 Locked ← OK →	5 Page 2	→ OK → Locked
6 ↓	6	↓
7 Unlock row 1	↑ row numbers	Unlocked row 2
8		

A.S.) Field level

It allows concurrent trans. to access the same ~~row~~ row as long as they required the use of different field [attributes] within the row.

Field level locking clearly displays the most flexible multi-user data access.

- The lock used by specified field can not be accessed by other trans. until the first one unlocks the field [attribute].

## B) Lock types.

- DBMS uses 2 different types of lock:
  - 1) Binary lock
  - 2) Shared / exclusive lock.

### B.1.) Binary lock.

- A binary lock has only 2 states: locked [1] or unlock [0].
- If an object such as database, table, page or row is locked by a trans. then no other trans. can use that object.
- If an object is unlocked, any trans. can lock the object for its use.
- A trans. must unlock the object after its termination.
- So, every trans. requires a lock and unlock operation for each accessed data items.
- The lock, unlock features eliminates the problem of lost update because lock is not realised until the WRITE statement is not completed.
- Binary locks are too restrictive.



Q.1) In the following example trans.  $T_1$  locks the PRODUCT table. So, trans.  $T_2$  can not access the same table until trans.  $T_1$  unlocks PRODUCT table.

Time	Trans.	Step	Stored value
1	$T_1$	Lock PRODUCT	
2	$T_1$	Read QTY	15
3	$T_1$	$QTY = QTY + 10$	
4	$T_1$	Write QTY	25
5	$T_1$	Unlock PRODUCT	
6	$T_2$	Lock PRODUCT	
7	$T_2$	Read QTY	25
8	$T_2$	$QTY = QTY - 10$	
9	$T_2$	Write QTY	15
10	$T_2$	Unlock PRODUCT	

Q.2) Shared / exclusive lock.

→ The shared lock exists when concurrent trans. are granted read access on the bases of a common lock.

→ A shared lock produces no conflict because the concurrent trans. are read only.

→ A shared lock issued when a trans. want to read data from a database and no exclusive lock is held on that data item.

Ex. If trans.  $T_1$  has a shared lock on data item  $X$  and trans.  $T_2$  wants to read data item  $X$  then  $T_2$  may also obtain a shared lock on data item  $X$ .

Shared lock		
Trans.	Operation	Shared lock
$T_1$	Read $X$	Allow
$T_2$	Read $X$	Allow

Exclusive lock.

- An exclusive lock is issued when a trans. want to update [write] a data item and no locks are currently held on that data item by any other trans.
- The exclusive lock is granted if and only if no other locks are held on the data item.
- This condition is called exclusive rule:- only one trans. at a time can own an exclusive lock on the object.

Ex. If trans.  $T_1$  updates the data item  $X$ , an exclusive lock is required by  $T_1$  over data item  $X$ .

Exclusive lock		
Trans.	Operation	Exclusive lock
$T_1$	Update $X$	Allow
$T_2$	Update $X$	Wait



- Using the shared / exclusive lock, a lock can have 3 states:  
Unlock, Shared [read], exclusive [Update, write]
- Shared / exclusive locks can lead to 2 major problems:
  - The resulting trans. schedule might not be serializable.
  - The schedule might create deadlock.
- A deadlock occurs when 2 trans. wait indefinitely for each other to unlock data.

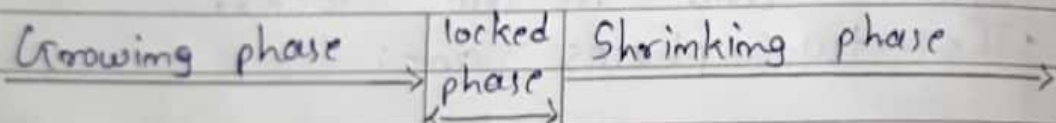
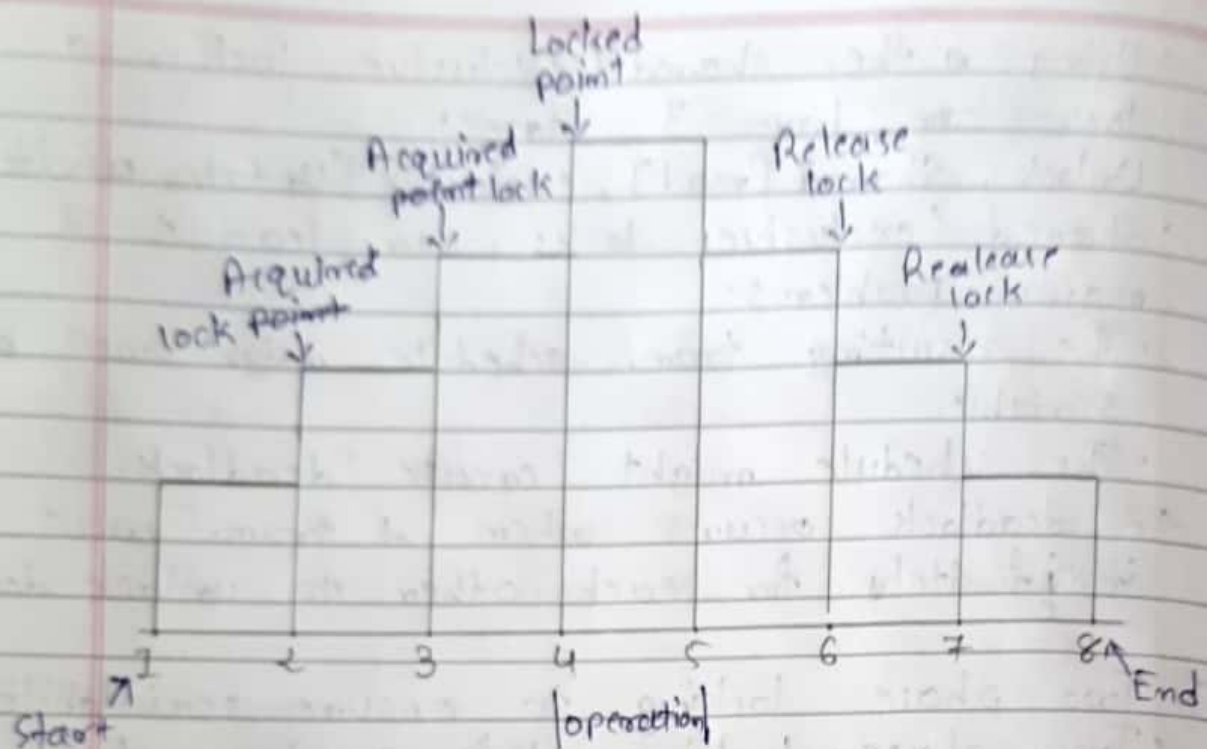
- c) Two phase locking to ensure serializability.
- Two phase locking defines how trans. acquires and release locks.
- The two phases are:

#### (i) Growing phase.

- In this phase, a trans. acquire all required locks without ~~an~~ unlocking data.
- Once all locks have been acquired, the trans. is in its locked point.

#### (ii) Shrinking phase.

- In this phase, a trans. release all locks and can not obtain a new lock.



✓ The two phase locking protocol has following rules:

- 1) Two trans. can not have conflicting locks.
- 2) No unlock operation can precede a lock operation.
- 3) No data are affected until the trans. is in its locked point.

✓ Two phase locking increase the trans. processing cost and might cause undesirable effects like deadlock.



22<sup>nd</sup> Feb. 2022  
Tuesday.

## 2) Deadlock.

• A deadlock occurs when 2 trans. wait indefinitely for each other to unlock data.

• For e.g., a deadlock occurs when 2 trans.  $T_1$  and  $T_2$  exist in the following mode:

$T_1$  = Access data item  $X$  &  $Y$

$T_2$  = Access data item  $Y$  &  $X$

If,  $T_1$  has not unlocked data item  $Y$ ,  $T_2$  can not begin.

If  $T_2$  has not unlocked data item  $X$ ,  $T_1$  can not begin.

$T_1$  and  $T_2$  both wait for the other to unlock the required data item.

• There are 3 basic techniques to control deadlock:

### A) Deadlock prevention.

- A trans. requesting a lock is aborted when there is a possibility that a deadlock can occur.
- If the trans. is aborted, all changes made by the trans. are rollback.

### B) Deadlock detection

- The DBMS periodically test the database for deadlock.
- If a deadlock is found, the victim trans. is aborted and other trans. continues.

c) Deadlock avoidance.

- The trans. must obtain all the locks it needs before it can be executed. This technique avoids conflicting trans.

\* Concurrency control with timestamping methods.

- The timestamping approach to scheduling concurrency trans. assigns a global, unique timestamp to each trans.
- Timestamps must have 2 properties:
  - Uniqueness  
It ensures that no equal timestamp value can exist.
  - Monotonicity.  
It ensures that timestamp value always increase.
- All database operations within the same trans. must have same timestamp.
- If 2 trans. conflict then one is stopped, rollback, reschedule and assign a new timestamp value.

Method:

Wait / Die & Wound / Wait scheme



## 2) Wait / Die

- Using wait / die scheme, older trans. waits for the younger one to complete and release its lock.
- If the trans. requesting the lock is older of 2 trans., it will wait until other trans. is completed and locks are released.
- If the trans. requesting the lock is younger of 2 trans., it will die [rollback] and is rescheduled using the same timestamp.

Trans. requesting lock	Trans. owning lock	Wait / Die scheme
$T_1(101)$	$T_2(102)$	$T_1$ waits until $T_2$ is completed & $T_2$ releases its locks.
$T_2(102)$	$T_1(101)$	$T_2$ dies [rollback], $T_2$ is rescheduled using same timestamp.

## 2) Wound/Wait

- In wound/wait scheme the older trans. rolls back the younger trans. and reschedule it.
- If the trans. requesting the lock is older of 2 trans., it will preempt [rollback/wound] the younger trans. and the younger trans. is rescheduled using the same timestamp.
- If the trans. requesting the lock is younger of 2 trans., it will wait until other trans. is completed and its locks are released.

Trans. requesting lock	Trans. owning lock	Wound/Wait scheme
$T_1(101)$	$T_2(102)$	$T_1$ preempts [rollback] $T_2$ . $T_2$ is rescheduled using same timestamp.
$T_2(102)$	$T_1(101)$	$T_2$ waits until $T_1$ is completed & $T_1$ releases its locks.



## Concurrency control with optimistic methods.

- The optimistic approach is based on the assumption that the database operation do not conflict.

- It requires neither locking nor timestamping methods.

- A trans. is executed without restrictions until it is committed.

- Using optimistic approach, each trans. has 3 phases: read, validation, write.

### A) Read phase.

- During read phase, the trans. reads the database, executes needed computation and makes the update to a private copy of the database value.

- All update operations of the trans. are recorded in a temporary update file.

### B) Validation phase.

- During this phase, the trans. is validated to ensure that the changes made will not affect integrity and consistency of the database.

- If the validation test is positive, the trans. goes to the write phase.

- If the validation test is negative, the trans. is restarted and the changes are discarded.

c) Write phase.

- During this phase, the changes are permanently applied to the database.

2<sup>nd</sup> Feb. 2022  
Wednesday.

\* Database recovery management.

- Database recovery restores a database from an inconsistent state to the consistent state.
- If a transaction operation cannot be completed for some reason, the transaction must be aborted and the changes are rollback.
- Critical events makes a database to stop working.
- This events can be
  - Hardware / software failure  
A failure of this type could be a hard disk failure, motherboard failure or application program and operating system failure.
  - Human caused incident  
This type of events can be categorised as unintentional or intentional.  
An unintentional failure is caused by a careless and user.  
Intentional events are more severe nature and indicate that the company data are at high risk.



- Natural disaster.

This category includes earthquake, fire, flood and power failure.

## → Transaction recovery.

Trans. recovery uses data in the trans. log to recover a database from inconsistent state to consistent state.

4 important concepts that affect the trans. recovery are

### A) Write-ahead-log protocol.

It ensures that trans. logs are always written before any database data are actually updated.

It also ensures that in case of a failure, the database can be recovered to a consistent state.

### B) Redundant-trans.-log

It ensures that physical disk failure will not impact the DBMS ability to recover data.

### C) Database buffer.

They are the temporary storage area in primary memory used to speed up disk operations.

### D) Database checkpoint

They are the operations in which DBMS write all the updated buffers to disk.

It is also registered in the trans. log.

→ Trans. recovery techniques.

→ Trans. recovery procedures make use of following 2 methods:

A) Deferred-write technique

B) Write-through technique

A) Deferred-write technique

→ When the recovery procedure uses deferred-write technique, the trans. operations do not immediately update the physical database.

→ Only trans. log is updated.

→ The database is physically updated only after the trans. reaches its commit point using information from the trans. log.

→ It is also called deferred update.

B) Write-through technique

→ When the recovery procedure uses a write-through technique, the database is immediately updated by trans. operation during execution, even before the commit point.

→ If the trans. aborts before its commit point, a ROLLBACK operation needs to be done to restore the database to a consistent state.

→ It is also called immediate update.



## Unit-2. Questions.

1. What is trans. ? Explain trans. property.  
[ACIDS property]
2. What is concurrency control? List problems associated with it. Explain any 1 in detail.
3. Explain lock granularity. [Levels of lock]
4. List and explain types of locks.
5. Explain 2 phase locking.
6. What is a deadlock? Explain its methods to control deadlock.
7. Explain concurrency control with timestamping method.  
or  
Explain wait/die and wound/wait method.
8. List and explain optimistic methods of concurrency control.
9. Explain different trans. recovery concepts with recovery techniques.